
TECH I CAN

Glytch

Curious today, Confident tomorrow

Glytch: A Beginner-Friendly Guide to AI Workflow Judgement

Classroom Edition (2026)

Tech I Can Series

By Paul McMurray

Understand. Check. Review. Decide.

Tech I Can Glytch

Curious today, Confident tomorrow

Glytch: A Beginner-Friendly Guide to AI Workflow Judgement

Paul McMurray

Classroom Edition (2026)

Dedicated to all of the budding techies of the future.
The world is ready for your brilliance.

Copyright

Copyright © 2026 Paul McMurray. This project is distributed under the MIT License. Permission is granted to use, copy, modify, merge, publish, distribute, sublicense, and sell copies, subject to inclusion of the copyright and license notice.

Published by Tech I Can, for practical AI literacy in schools and workshops. See the repository LICENSE file for full terms and warranty disclaimer.

Publication Data

Publisher: Tech I Can

Publication date: May 2026

ISBN: Not assigned (Edition 1 school release, not for retail sale)

Edition: Classroom Edition (2026)

Trim size: A4 digital classroom edition

Tech I Can

Curious today, Confident tomorrow.

Printed in digital format for classroom distribution.

For classroom and school use, the Disclaimer on the following page applies alongside the licence terms above.

Disclaimer

Glytch is a workshop-safe simulation for education and discussion.

It is not a production automation platform, not a decision engine for high-stakes contexts, and not a replacement for teacher, manager, or professional judgement.

Always apply local safeguarding, data protection, and policy requirements in school and workplace environments.

About the Author

Teaching intent: The teaching intent is simple: clear steps, honest evidence, and confidence through practice.

I'm Paul McMurray, a software engineer, problem-solver, and lifelong learner from the North East of England.

My work has taken me deep into the world of technology, reliability, automation, and complex systems, but the part I enjoy most is making difficult things easier to understand. Whether I'm building tools, improving systems, writing guidance, or helping others find a clearer way forward, I'm drawn to practical ideas that make a real difference.

Outside of work, I'm a dad, a music lover, and someone who believes in staying curious. I enjoy soul, funk, and disco, long walks, good food, and finding new ways to teach my son how the world works - from computers and games to kindness, resilience, and asking good questions.

This book is personal to me because it reflects how I think: take something complicated, break it down, make it useful, and leave people with something they can actually do.

I don't pretend to have all the answers, but I've learned that progress often starts with asking better questions, being honest about what you don't know, and being willing to keep going.

Icons Used in This Book

You will see these symbols throughout the chapters. They signal what kind of guidance is being offered.

	Lightbulb Takeaway	A practical classroom takeaway you should remember.
	Definition	A clear meaning for a technical term used in the chapter.
	Note	Useful context, guardrails, or teaching detail for this step.
	Snippet Purpose	Why this command is being run before you use it.
	Snippet Change	What changed from a previous command and why it matters.

Contents

Chapter 1: AI Beyond The Chat Box	10
Chapter 2: The Lowest Safe Level Principle	13
Chapter 3: Evidence And Grounding	19
Chapter 4: Critique And Revision	22
Chapter 5: Controlled Loops And Limits	25
Chapter 6: Orchestration And Taskboards	28
Chapter 7: Classroom Delivery And Assessment	31
Chapter 8: Pilot Planning And Continuous Improvement	34

**Foundations**

Build confidence with the core workflow.

Chapter Overview

This chapter focuses on 'AI Beyond The Chat Box'. You will examine why polished output is not enough, how to separate fluency from reliability, and how Glytch levels support comparison.

What you will walk away with

- why polished output is not enough
- how to separate fluency from reliability
- how Glytch levels support comparison

Chapter Focus: Chapter 1: AI Beyond The Chat Box

Chapter 1: AI Beyond The Chat Box

About this chapter

This chapter sets the foundation for the whole book. You will frame AI as a capability journey, not a tool popularity contest. The goal is to make clear why one prompt can be useful for some tasks but risky for others.

What you are going to use

- Glytch running locally
- one safe scenario
- your observation notes

What you will learn in this chapter

- why polished output is not enough
- how to separate fluency from reliability
- how Glytch levels support comparison

The work, clearly laid out

1. choose one scenario
2. run Level 1 once
3. list what is useful and what is uncertain
4. define one human review checkpoint

Examples of what you might see

- clear wording but unsupported claims
- confident tone with missing caveats
- useful starting structure that still needs review

Why This Matters

People often trust style before they test quality. That is risky in classrooms, teams, and public services. Early awareness reduces over-trust and creates better habits.

Action 1: What You Learned

- A clean answer can still hide weak assumptions.
- Human review is not optional for meaningful tasks.
- AI literacy starts with process awareness.

Action 2: Reflect

- Which part of the Level 1 output would you verify first?
- What looked convincing but remained unproven?
- What should never be accepted without review in your context?

Action 3: Do This Next

- Pair with a colleague and compare your risk notes for the same output.
- Share one verification checklist item with your group.

**Foundations**

Build confidence with the core workflow.

Chapter Overview

This chapter focuses on 'The Lowest Safe Level Principle'. You will examine what changes between prompt and checked workflow, how to choose level based on risk, and why "more advanced" does not automatically mean "better".

What you will walk away with

- what changes between prompt and checked workflow
- how to choose level based on risk
- why "more advanced" does not automatically mean "better"

Chapter Focus: Chapter 2: The Lowest Safe Level Principle

Chapter 2: The Lowest Safe Level Principle

About this chapter

This chapter introduces the central decision rule in Glytch: use the lowest level that safely gets the task done. You will learn why higher complexity is sometimes justified and sometimes wasteful.

What you are going to use

- one repeated scenario
- Level 1 and Level 3 runs
- a side-by-side comparison table

What you will learn in this chapter

- what changes between prompt and checked workflow
- how to choose level based on risk
- why "more advanced" does not automatically mean "better"

The work, clearly laid out

1. run Level 1 for baseline
2. run Level 3 on the same prompt
3. compare claims, checks, and confidence
4. decide which level is sufficient

Examples of what you might see

- Level 3 adds clearer intermediate logic
- tool-assisted checks reduce obvious arithmetic errors
- baseline remains faster for low-risk drafting

The 8 levels, layer by layer

This section gives a practical understanding of what each level adds, what that layer is doing, and what users should understand before moving up.

Level 1 - Prompt baseline

What this layer does:

- Takes one prompt and returns one draft response.
- Keeps human decision-making fully external.

What users should understand:

- Fast drafting value is high.
- Reliability and traceability are low unless the user adds checks manually.

Best use:

- Early brainstorming, wording ideas, low-risk first drafts.

Level 2 - Instruction framing

What this layer does:

- Adds stronger instructions, constraints, and output format expectations.
- Reduces ambiguity compared with raw prompting.

What users should understand:

- Better structure does not guarantee correctness.
- Compliance with format can hide weak reasoning if evidence is not checked.

Best use:

- Structured summaries, checklists, first-pass procedural outputs.

Level 3 - Tool and check integration

What this layer does:

- Uses bounded helper tools and checkable intermediate steps.
- Makes parts of the run more inspectable than pure prompting.

What users should understand:

- Tool output can improve consistency but still needs interpretation.

- Human review is still required for edge cases and context.

Best use:

- Tasks where small calculation or retrieval checks reduce avoidable mistakes.

Level 4 - Evidence grounding

What this layer does:

- Links claims to supplied source material.
- Encourages answer construction from known context instead of free-form guessing.

What users should understand:

- Grounding improves auditability.
- Missing or weak source material still leads to weak outputs.

Best use:

- Fact-sensitive educational, operational, or policy discussions.

Level 5 - Planning and decomposition

What this layer does:

- Breaks broad tasks into explicit sub-steps.
- Makes process intent visible before final output.

What users should understand:

- Better planning can reduce confusion.
- Excess planning can add overhead for simple tasks.

Best use:

- Multi-step tasks where order and scope matter.

Level 6 - Critique and revision

What this layer does:

- Produces draft, critique, and improved draft flow.
- Separates initial generation from quality improvement.

What users should understand:

- Revision improves clarity and robustness when critique criteria are strong.
- Polished wording after revision is still not proof of truth.

Best use:

- Deliverables that need quality uplift before human approval.

Level 7 - Controlled loop

What this layer does:

- Repeats observe-act-check style cycles with boundaries.
- Adds explicit stop and control mechanics.

What users should understand:

- Loops increase process power and process risk.
- Stop conditions and escalation criteria must be clear.

Best use:

- Iterative tasks where each pass should be reviewed against a goal.

Level 8 - Orchestration

What this layer does:

- Coordinates specialist worker roles with verifier and approval gates.
- Produces taskboard, merge decision path, and final status signal.

What users should understand:

- Highest structure means highest governance burden.
- A merged status is a process result, not automatic factual validity.

Best use:

- Complex tasks where coordination, traceability, and explicit human approvals are essential.

Quick comparison table

Level	Primary value	Main risk if misused
1	speed and ideation	over-trust in fluent draft text

2	clearer structure	false confidence from tidy formatting
3	checkable helper steps	tool output accepted without context review
4	evidence traceability	weak source quality mistaken for strong evidence
5	process clarity	unnecessary complexity for simple work
6	quality improvement loop	polish mistaken for correctness
7	bounded iteration	control drift if stop rules are weak
8	coordinated governance	process confidence replacing human judgement

Why This Matters

Teams waste time when they over-engineer small tasks and create risk when they under-control important tasks. The lowest safe level principle balances both.

Action 1: What You Learned

- Level choice should follow task risk, not feature excitement.
- Checks add value when mistakes are costly.
- Simpler flow can be the right answer.

Action 2: Reflect

- Which scenario in your school or team is over-complicated today?
- Which scenario is under-checked and needs more control?
- How would you explain "lowest safe level" to a new colleague?

Action 3: Do This Next

- Work in pairs to classify three tasks by appropriate level.
- Compare your classifications with another pair and resolve disagreements.

**Foundations**

Build confidence with the core workflow.

Chapter Overview

This chapter focuses on 'Evidence And Grounding'. You will examine what grounding looks like in output, how to inspect whether claims map to supplied material, and where evidence still leaves judgement questions open.

What you will walk away with

- what grounding looks like in output
- how to inspect whether claims map to supplied material
- where evidence still leaves judgement questions open

Chapter Focus: Chapter 3: Evidence And Grounding

Chapter 3: Evidence And Grounding

About this chapter

This chapter focuses on Level 4 behavior and the difference between unsupported claims and evidence-tied claims. You will practice tracing claims to provided material so learners can see what is substantiated, what is uncertain, and what should be revised before use.

What you are going to use

- Level 4 run
- teacher-provided or facilitator-provided source notes
- a claim-check table

What you will learn in this chapter

- what grounding looks like in output
- how to inspect whether claims map to supplied material
- where evidence still leaves judgement questions open

The work, clearly laid out

1. run a scenario with evidence notes
2. highlight claims in the output
3. map each claim to provided evidence
4. flag unsupported statements

Examples of what you might see

- direct claims that clearly map to source material
- generic claims not supported by provided notes
- mixed outputs where only part of the answer is grounded

Why This Matters

Evidence discipline is essential in education, operations, and public communication. Grounding creates traceability and reduces accidental misinformation.

Action 1: What You Learned

- Evidence-backed statements are easier to trust and audit.
- Unsupported claims should be treated as draft material.
- Grounding reduces, but does not eliminate, risk.

Action 2: Reflect

- Which claim from your run had the strongest evidence link?
- Which claim should be removed until better support exists?
- What evidence standard should your team adopt?

Action 3: Do This Next

- Swap outputs with a partner and run independent claim checks.
- Build a shared evidence checklist with your group.

**Foundations**

Build confidence with the core workflow.

Chapter Overview

This chapter focuses on 'Critique And Revision'. You will examine why first draft quality is only a starting point, how critique changes reliability and clarity, and when revision loops should stop.

What you will walk away with

- why first draft quality is only a starting point
- how critique changes reliability and clarity
- when revision loops should stop

Chapter Focus: Chapter 4: Critique And Revision

Chapter 4: Critique And Revision

About this chapter

This chapter explores Level 6 patterns where the first draft is reviewed and improved before wider use. You will look at revision quality in practical terms: clearer structure, stronger caveats, better reasoning flow, and reduced risk language before any release decision.

What you are going to use

- Level 6 output
- critique notes
- revised output comparison

What you will learn in this chapter

- why first draft quality is only a starting point
- how critique changes reliability and clarity
- when revision loops should stop

The work, clearly laid out

1. run Level 6 once
2. identify weak points in the first draft
3. review revised output
4. assess if revision addressed the weaknesses

Examples of what you might see

- improved structure after critique
- clearer caveats in revised text
- unresolved risk despite stylistic polish

Why This Matters

Many AI failures happen because first drafts are treated as final. Revision introduces accountability and reduces avoidable errors.

Action 1: What You Learned

- Revision is a quality process, not an optional extra.
- Strong critique criteria produce better second drafts.
- Not every revision cycle improves truthfulness.

Action 2: Reflect

- Which critique point produced the biggest improvement?
- What risk remained after revision?
- When should a human decide to stop revising and escalate?

Action 3: Do This Next

- Pair with a classmate to critique each other's run outputs.
- Agree one common revision rubric for your team.

**Foundations**

Build confidence with the core workflow.

Chapter Overview

This chapter focuses on 'Controlled Loops And Limits'. You will examine why loops can help and harm, how to interpret loop status safely, and what stop rules should be explicit.

What you will walk away with

- why loops can help and harm
- how to interpret loop status safely
- what stop rules should be explicit

Chapter Focus: Chapter 5: Controlled Loops And Limits

Chapter 5: Controlled Loops And Limits

About this chapter

This chapter introduces loop-based behavior and control boundaries. You will focus on how repeated AI actions need explicit stop conditions and monitoring.

What you are going to use

- Level 7 run
- loop summary
- stop-condition checklist

What you will learn in this chapter

- why loops can help and harm
- how to interpret loop status safely
- what stop rules should be explicit

The work, clearly laid out

1. run Level 7
2. inspect loop phases
3. identify stop and check conditions
4. record one potential runaway risk

Examples of what you might see

- clear observe-act-check progression
- incomplete checks in one iteration
- loop output that still needs human approval

Why This Matters

Without clear boundaries, iterative systems can amplify mistakes quickly. Responsible loop design improves safety and predictability.

Action 1: What You Learned

- Loops need clear limits and accountability.
- Repetition without verification can increase error propagation.
- Human checkpoints remain essential.

Action 2: Reflect

- Which stop condition felt most important in your run?
- What signal would tell you to halt the loop?
- How would you explain loop safety to a beginner group?

Action 3: Do This Next

- In small groups, draft a loop stop-policy card for one scenario.
- Compare cards and agree the strongest boundary language.

**Foundations**

Build confidence with the core workflow.

Chapter Overview

This chapter focuses on 'Orchestration And Taskboards'. You will examine how worker decomposition changes workflow, what verifier and approval gates signal, and why "merged" is not the same as "perfect".

What you will walk away with

- how worker decomposition changes workflow
- what verifier and approval gates signal
- why "merged" is not the same as "perfect"

Chapter Focus: Chapter 6: Orchestration And Taskboards

Chapter 6: Orchestration And Taskboards

About this chapter

This chapter centers on Level 8, where multiple worker roles coordinate through verification and approval gates. You will focus on interpreting orchestration signals responsibly so teams can avoid false confidence and keep named human accountability at final decision points.

What you are going to use

- Level 8 run
- taskboard view
- approval and merge summary

What you will learn in this chapter

- how worker decomposition changes workflow
- what verifier and approval gates signal
- why "merged" is not the same as "perfect"

The work, clearly laid out

1. run Level 8 with your scenario
2. inspect each worker output summary
3. review verifier and approval fields
4. decide whether you would release or revise

Examples of what you might see

- one worker blocked while others complete
- verifier warning before merge
- merge deferred pending human review

Why This Matters

Orchestration adds control possibilities and new failure points. Understanding both is critical for responsible deployment decisions.

Action 1: What You Learned

- Coordination does not remove review obligations.
- Merge outcomes are process indicators, not truth guarantees.
- Worker specialization can improve clarity when roles are explicit.

Action 2: Reflect

- Which worker role added the most value in your run?
- Where could orchestration create false confidence?
- What should always require a named human approver?

Action 3: Do This Next

- Review another team's taskboard and challenge one merge decision.
- Co-design a two-person approval workflow with a partner.

**Foundations**

Build confidence with the core workflow.

Chapter Overview

This chapter focuses on 'Classroom Delivery And Assessment'. You will examine how to sequence levels in limited time, which misconceptions to address early, and how to evaluate understanding without overloading learners.

What you will walk away with

- how to sequence levels in limited time
- which misconceptions to address early
- how to evaluate understanding without overloading learners

Chapter Focus: Chapter 7: Classroom Delivery And Assessment

Chapter 7: Classroom Delivery And Assessment

About this chapter

This chapter helps educators and facilitators run effective sessions, keep safety boundaries visible, and collect useful learning evidence. You will connect classroom design decisions to practical outcomes so each session produces better discussion, clearer reflection, and safer AI use habits.

What you are going to use

- teacher guide
- student worksheet
- one session reflection template

What you will learn in this chapter

- how to sequence levels in limited time
- which misconceptions to address early
- how to evaluate understanding without overloading learners

The work, clearly laid out

1. choose a 20, 45, or 90 minute format
2. select one safe scenario
3. run levels with comparison prompts
4. collect learner reflections and level justifications

Examples of what you might see

- learners equating confidence with correctness
- learners over-valuing higher levels
- stronger reflection quality after side-by-side comparisons

Why This Matters

AI literacy succeeds when learners can justify decisions, not when they memorize terms. Structured reflection builds transferable judgement.

Action 1: What You Learned

- Time-boxed comparison is often more effective than feature tours.
- Misconceptions are predictable and teachable.
- Reflection prompts can surface true understanding.

Action 2: Reflect

- Which prompt generated the best learner reasoning?
- What misconception appeared most often in your context?
- How can you make safety boundaries more visible next session?

Action 3: Do This Next

- Run a paired facilitation rehearsal with a colleague.
- Exchange student worksheet feedback with another teacher.

**Foundations**

Build confidence with the core workflow.

Chapter Overview

This chapter focuses on 'Pilot Planning And Continuous Improvement'. You will examine how to start small and measure clearly, what evidence should guide iteration, and how to decide readiness for wider rollout.

What you will walk away with

- how to start small and measure clearly
- what evidence should guide iteration
- how to decide readiness for wider rollout

Chapter Focus: Chapter 8: Pilot Planning And Continuous Improvement

Chapter 8: Pilot Planning And Continuous Improvement

About this chapter

This chapter explains how to run a small pilot, gather meaningful feedback, and improve your rollout responsibly. You will use pilot evidence to decide what should scale, what should change, and what needs stronger controls before broader adoption.

What you are going to use

- pilot plan template
- before-and-after confidence prompts
- release gate checklist

What you will learn in this chapter

- how to start small and measure clearly
- what evidence should guide iteration
- how to decide readiness for wider rollout

The work, clearly laid out

1. choose one cohort
2. run one baseline session
3. collect confidence and clarity signals
4. review gaps and revise materials

Examples of what you might see

- confidence increase but uncertainty about evidence standards
- strong staff engagement with workflow detail
- need for clearer guidance on human approval gates

Why This Matters

Pilots reduce rollout risk, reveal practical teaching gaps, and create evidence for informed scaling decisions.

Action 1: What You Learned

- Small pilots create higher-quality improvements than big launches.
- Feedback should include both confidence and caution signals.
- Release gates prevent quality drift.

Action 2: Reflect

- What one indicator best shows your pilot helped judgement?
- Which uncertainty remained strongest after the session?
- What should be improved before expanding to more groups?

Action 3: Do This Next

- Debrief pilot findings with a partner team.
- Co-author one improvement sprint plan with clear owners.

References (APA 7th Edition)

Bender, E. M., Gebru, T., McMillan-Major, A., & Shmitchell, S. (2021). On the dangers of stochastic parrots: Can language models be too big? Proceedings of the 2021 ACM Conference on Fairness, Accountability, and Transparency, 610-623.

Floridi, L., & Cowls, J. (2019). A unified framework of five principles for AI in society. Harvard Data Science Review, 1(1).

Ng, A. (2021). Machine learning yearning (Draft).

OpenAI. (2024). Usage policies. <https://openai.com/policies>

UNESCO. (2021). Recommendation on the ethics of artificial intelligence. UNESCO Publishing.

One-line close

Glytch helps people build practical AI judgement by making workflow control, evidence, and human review visible.